

On the (Statistical) Detection of Adversarial Examples

Kathrin Grosse[†], Praveen Manoharan[‡], Nicolas Papernot[‡], Michael Backes^{†*}, Patrick McDaniel[‡]

CISPA, Saarland Informatics Campus[†]; Penn State University[‡]; MPI SWS*
kathrin.grosse@cispa.saarland

ABSTRACT

Machine Learning (ML) models are applied in a variety of tasks such as network intrusion detection or malware classification. Yet, these models are vulnerable to a class of malicious inputs known as adversarial examples. These are slightly perturbed inputs that are classified incorrectly by the ML model. The mitigation of these adversarial inputs remains an open problem.

As a step towards understanding adversarial examples, we show that they are **not drawn from the same distribution than the original data**, and can thus be **detected using statistical tests**. Using this knowledge, we introduce a **complimentary approach to identify specific inputs that are adversarial**. Specifically, we augment our ML model with an additional output, in which the model is trained to classify all adversarial inputs.

We evaluate our approach¹ on multiple adversarial example crafting methods (including the fast gradient sign and saliency map methods) with several datasets. The statistical test flags sample sets containing adversarial inputs confidently at sample sizes between 10 and 100 data points. Furthermore, our augmented model either detects adversarial examples as outliers with high accuracy (> 80%) or increases the adversary’s cost—the perturbation added—by more than 150%. In this way, we show that statistical properties of adversarial examples are essential to their detection.²

1 INTRODUCTION

Machine learning algorithms are usually designed under the assumption that models are trained on **samples drawn from a distribution that is representative of test samples** for which they will later make predictions—ideally, the training and test distributions should be identical. However, this does not hold in the presence of adversaries. A motivated adversary may either manipulate the training [4] or test [33] distribution of a ML system. This has severe consequences when ML is applied to security-critical problems. Attacks are increasingly elaborate, as demonstrated by the variety of strategies available to evade malware detection built with ML [33, 37].

Often, adversaries construct their attack inputs from a benign ML input. For instance, the feature vector of a malware—correctly classified by a ML model as malware—can be modified into a new feature vector, the *adversarial example*, that is classified as benign [14, 37].

Defenses proposed to mitigate adversarial examples, such as adversarial training [11] and defensive distillation [29], all fail to adapt to changes in the attack strategy. They both make it harder for the adversary to craft adversarial examples using existing techniques

only, thus creating an arms race [7, 25]. However, we argue that this arms race is not inevitable: by definition, adversarial examples must exhibit some statistical differences with the legitimate data on which ML models perform well.

Hence, we develop in this work a countermeasure that uses the distinguishability of adversarial examples with respect to the expected data distribution. We use statistical testing to evaluate the hypothesis that adversarial examples, crafted to evade a ML model, are outside of the training distribution. We show that the hypothesis holds on diverse datasets [1, 2, 19] and adversarial example algorithms [11, 26, 27].

However, this test needs to be presented with a sufficiently large sample set of suspicious inputs—as its confidence diminishes with the number of malicious inputs in the sample set. Therefore, we propose a second complimentary mechanism for detecting individual adversarial examples. The idea also exploits the statistical distinguishability of adversarial examples to design an outlier detection system, but this time it is directly integrated in the ML model. Indeed, we show that models can be augmented with an additional output reserved for adversarial examples—in essence training the model with adversarial examples as their own class. The model, trained to map all adversarial inputs to the added output, exhibits robustness to adversaries.

Our contributions are the following:

Statistical Test—In Section 5, we employ a statistical test to distinguish adversarial examples from the model’s training data. Among tests proposed in the literature, we select the kernel-based two-sample test introduced by Gretton et al. [13]. This test has the key benefit of being model-agnostic; because its kernel allows us to apply the test directly on samples from the ML model’s input data.

We demonstrate the good performance of this test on three datasets: MNIST (hand-written digits), DREBIN (Android malware) and MicroRNA (medical data). Specifically, we show that the test can confidently detect samples of 50 adversarial inputs when they differ from the expected dataset distribution. Results are consistent across multiple generation techniques for adversarial examples, including the fast gradient sign method [11] and the Jacobian-based saliency map approach [27].

Integrated Outlier Detection—As the statistical test’s confidence diminishes when it is presented with increasingly small sample sets of adversarial inputs, we propose another outlier detection system. We add an additional class to the model’s output, and train the model to recognize adversarial examples as part of this new class. The intuition behind the idea is the same (detect adversarial examples using their statistical properties) but this approach allows the defender to detect individual adversarial examples among a set of inputs identified as malicious (by the statistical test for instance).

¹Please contact the authors to obtain the code for reproduction of the experiments.

²Recent work [8], however, has shown that our approach is vulnerable if optimization-based attacks are used, which require however more computational effort.

We observe here that adversarial examples lie in unexpected regions of the model’s output surface because they are not representative of the distribution. By training the model to identify out-of-distribution inputs, one removes at least part of its error away by filling in the things that are demonstrable (e.g., adversarial examples).

In Section 6, we find that this approach correctly assigns adversarial examples as being part of the outlier class with over 80% success for two of the three datasets considered. For the third dataset, they are not frequently detected but the perturbation that an adversary needs to add to mislead the model is increased by 150%. Thus, the cost of conducting an attack is increased in all cases. In addition, adversarial examples that are not detected as outliers because they are crafted with small perturbations are often correctly classified by our augmented model: the class of the legitimate input from which they were generated is *recovered*.

Arms race—We then investigate adversarial strategies taking into account the defense deployed. For instance, black-box attacks were previously shown to evade adversarial training and defensive distillation [25]. The adversary uses an auxiliary model to find adversarial inputs that are also misclassified by the defended model (because the defended model makes adversarial crafting harder but does not solve the model error). Our mechanisms perform well under such black-box scenarios: adversarial inputs crafted by a black-box attack are more likely to be detected than those computed directly by an adversary with access to our model.

2 BACKGROUND

We provide here the relevant background on ML and adversarial ML. We finally give an overview of the **statistical hypothesis test applied in this paper**.

2.1 Machine Learning Classifiers

We introduce ML notation used throughout this paper. All ML models considered are classifiers and learn a function $f(x) \mapsto y$. An input point or example $x \in X$ is made up of n components or features (e.g., all system level calls made by an Android application), and $y \in Y$ is a label (e.g., malware or benign). In classification problems, the possible values of y are discrete. The output of the model, however, is often real valued probabilities over the set of possible labels, from which the most likely label is inferred as the one with the largest probability.

In other words, there is an underlying and almost always unknown distribution $D_{\text{real}}^{C_i}$ for each class C_i . The set of training data X is sampled from this distribution, and the classifier approximates this distribution during training, thereby learning $D_{\text{train}}^{C_i}$. The set of test data X_t , used to validate the classifier’s performance, is assumed to be drawn from the same $D_{\text{real}}^{C_i}$.

Next, we present typical ML models used to solve classification problems and studied in this paper.

Decision Trees—These models are composed of internal nodes and leafs, whose graph makes up a tree. Each leaf is assigned one of the possible labels, while the intermediate nodes form a path of conditions defined using the input features. An example is classified by finding a path of appropriate conditions from the root to one of the leaves. Decision trees are created by successively maximizing

the information gain resulting from the choice of a condition as a way to partition the data in two subsets (according to the value of an input feature).

Support Vector Machines—They compute a $n - 1$ dimensional hyperplane to separate the training points. Since there are infinitely many such hyperplanes, the one with the largest margin is computed—yielding a convex optimization problem given the training data (X, Y) .

Neural Networks—They are composed of small computational units called *neurons* that apply an *activation function* to their weighted input. Neurons are organized in interconnected layers. Depending on the number of layers, a network is said to be *shallow* (single intermediate layer) or *deep* (several intermediate layers). Information is propagated through the network by having the output of a given layer be the input of the following layer. Each of these links is parameterized by weights. The set of model weights—or model parameters—are trained to minimize the model’s prediction error $\|f(x) - y\|$ on a collection of known input-output pairs (X, Y) .

Logistic regression—This linear model can be conceptualized as a special case of neural networks without hidden layers. For problems with two classes, the logistic function is the activation function. For multi-class problems, it is the softmax. They are trained like neural nets.

2.2 Adversarial Machine Learning

Adversarial ML [17], and more generally the security and privacy of ML [28], is concerned with the study of vulnerabilities that arise when ML is deployed in the presence of malicious individuals. Different attack vectors are available to adversaries. They can target ML during training [4] of the model parameters or during test time [33] when making predictions.

In this paper, we defend against test time attacks. They target a trained model $f(_, \theta)$, and typically aim to find an example x' similar to an original example x , which is however classified differently. To achieve this, a perturbation δ with same dimensionality as x is computed:

$$f(x', \theta) \neq f(x, \theta) \text{ where } x' = x + \delta \text{ and } \min \delta$$

where δ is chosen to be minimal to prevent detection and as to indirectly represent the attackers limitations when perturbing features. When targeting computer vision, the perturbation must not be detectable to the human eye. When targeting a malware detector, the perturbation must not remove the application’s malicious behavior. Instead of simply having inputs classified in a wrong class, the attacker can also target a particular class.

A typical example of such attacks is the evasion of a bayesian spam filter, first demonstrated by Lowd et al. [22]. Malware detection systems have also been targeted, as shown by Srndic et al. [33] or Grosse et al. [14]. In addition, these adversarial inputs are known to transfer across (i.e., to mislead) multiple models simultaneously [34]. This transferability property was used to create attacks against black-box ML systems in settings where the adversary has no access to the model or training data [25, 26]. A detailed discussion of some of these attacks can be found in Section 4.

Several defenses for attacks at test time have been proposed. For instance, training on adversarial inputs pro-actively [11] or

performing defensive distillation [29]. Both of them may fail due to gradient masking [25]. Other approaches make use of game theory [6, 9, 21]. However, they are computationally expensive.

2.3 Statistical Hypothesis Testing

The framework of **two-sample statistical hypothesis testing** was introduced to determine whether **two randomly drawn samples originate from the same distribution**.

Formally, let $X \sim p$ denote that sample X was drawn from a distribution p . A statistical test can then be formalized as follows: let $X_1 \sim p$, where $|X_1| = n$ and $X_2 \sim q$, where $|X_2| = m$. The null hypothesis H_0 states that $p = q$. The alternative hypothesis, H_A , on the other hand, is that $p \neq q$. The statistical test $\mathcal{T}(X_1, X_2) : \mathcal{X}^n \times \mathcal{X}^m \rightarrow \{0, 1\}$ takes both samples as its input and distinguishes between H_0 and H_A . In particular, the p-value returned is matched to a significance level, denoted α . **The p-value is the probability that we obtain the observed outcome or a more extreme one.** α relates to the confidence of the test, and an according threshold is fixed before the application of the test, **typically at 0.05 or 0.01**. If the p-value is smaller than the threshold, H_0 is rejected. A consistent test will reject H_0 when $p \neq q$ in the large sample size limit.

There are several two-sample tests for higher dimensions. For instance, **the Hotellings T^2 test evaluates whether two distributions have the same mean** [16]. Several other tests depending on graph or tree properties of the data were proposed by Friedman et al. [10], Rosenbaum et al. [31] or Hall et al. [15].

Most of these tests are not appropriate when considering data with high dimensionality. **This led Gretton et al. [13] to introduce a kernel-based test.** In this case, we measure the distance between two probabilities (represented by samples X_1 and X_2). **In practice, this distance is formalized as the biased estimator of the true Maximum Mean Discrepancy (MMD):**

$$\text{MMD}_b[\mathcal{F}, X_1, X_2] = \sup_{f \in \mathcal{F}} \left(\frac{1}{n} \sum_{i=1}^n f(x_{1i}) - \frac{1}{m} \sum_{i=1}^m f(x_{2i}) \right)$$

where the maximum indicates that we pick the kernel function f from the function class \mathcal{F} that maximize the difference between the functions. Further, in contrast to other measures, we do not need the explicit probabilities.

Gretton et al. [13] introduced several tests. We focus on one of them in the following: a test based on the asymptotic distribution of the unbiased MMD. However, to consistently estimate the distribution of the MMD under H_0 , we need to bootstrap³. Here, bootstrapping refers to a subsampling method, where one samples from the data available with replacement. By repeating this procedure many times, we obtain an estimate for the MMD value under H_0 .

3 METHODOLOGY

Here, we introduce a **threat model to characterize the adversaries our system is facing**. We also derive a formal argument justifying the statistical divergence of adversarial examples from benign training points. This observation underlies the design of our defensive mechanisms.

³Other methods have been proposed, such as moment matching Pearson curves. We focus here on one specific test used in this paper.

3.1 Threat model

Adversarial knowledge— Adversarial example crafting algorithms proposed in the literature primarily differ in the assumptions they make about the knowledge available to adversaries [28]. Algorithms fall in two classes of assumptions: *white-box* and *black-box*.

Adversaries operating in the white-box threat model have unfettered access to the ML system’s architecture, the value of its parameters, and its training data. In contrast, other adversaries do not have access to this information. They operate in a black-box threat model where they typically can interact with the model only through an interface analog to a cryptographic oracle: it returns the label or probability vector output by the model when presented with an input chosen by the adversary.

In this paper, we are designing a defensive mechanism. As such, we must consider the worst-case scenario of the strongest adversary. We therefore operate in both the white-box and the black-box threat model. While our attacks may not be practical for certain ML systems, it allows us to provide stronger defensive guarantees.

Adversarial capabilities— These are only restricted by constraints on the perturbations introduced to craft adversarial examples from legitimate inputs. Such constraints vary from dataset to dataset, and as such we leave their discussion to the description of our setup in Section 4.

3.2 Statistical Properties of Adversarial Examples

When learning a classifier from training data as described in Section 2, one seeks to learn the real distributions of features $D_{\text{real}}^{C_i}$ for each subset C_i corresponding to a class i . These subsets define a partition of the training data, i.e. $\cup_i C_i = X$. However, due to the limited number of training examples, any machine learning algorithm will only be able to learn an approximation of this real distribution, the *learned feature distributions* $D_{\text{train}}^{C_i}$.

A notable result in ML is that any *stable* learning algorithms will learn the real distribution $D_{\text{real}}^{C_i}$ up to any multiplicative factor given a sufficient number of training examples drawn from $D_{\text{real}}^{C_i}$ [5]. Stability refers here to the fact that given a slight modification of the data, **the resulting classifier and its prediction do not change much**. Coming back to our previous reasoning, however, this *full generalization* is in practice impossible due to the finite (and often small) number of training examples available.

The existence of adversarial examples **is a manifestation of the difference between the real feature distribution $D_{\text{real}}^{C_i}$ and the learned feature distribution $D_{\text{train}}^{C_i}$** : the adversary follows the strategy of **finding a sample drawn from $D_{\text{real}}^{C_i}$ that does not adhere to the learned distribution $D_{\text{train}}^{C_i}$** . This is only partially dependent on the actual algorithm used to compute the adversarial example. Yet, the adversary (or any entity as a matter of fact) does not know the real feature distribution $D_{\text{real}}^{C_i}$ (otherwise one could use that distribution in lieu of the ML model). Therefore, existing crafting algorithms generate adversarial examples by perturbing legitimate examples drawn from $D_{\text{train}}^{C_i}$, as discussed in Section 2.2.

Independently of how adversarial examples were generated, all adversarial examples for a class C_i will constitute a new distribution

$D_{\text{adv}}^{C_i}$ of this class. Following the above arguments, clearly $D_{\text{adv}}^{C_i}$ is consistent with $D_{\text{real}}^{C_i}$, since each adversarial example for a class C_i is still a data point that belongs to this class. On the other hand, however, $D_{\text{adv}}^{C_i} \neq D_{\text{train}}^{C_i}$. This follows from a reductio ad absurdum: if the opposite was true, adversarial examples would be correctly classified by the classifier.

As discussed in Section 2.3, consistent statistical tests can be used to detect whether two sets or samples X_1 and X_2 were sampled from the same distribution or not. A sufficient (possibly infinite) number of examples in each sample allows such a consistent statistical test to detect the difference in the distributions even if the underlying distributions of X_1 and X_2 are very similar.

Following from the above, statistical tests are natural candidates for adversarial example detection. Adversarial examples have to **inherently be distributed differently from legitimate examples used during training. The difference in distribution should consequently be detectable by a statistical test.** Hence, the **first hypothesis we want to validate or invalidate is the ability of a statistical test to distinguish between benign and adversarial data points.** We have two practical limitations, one is that we can do so by observing a finite (and small) number of examples, the second that we are restricted to existing adversarial example crafting algorithms.

HYPOTHESIS 1. We only need a bounded number of n examples to observe a measurable difference in the distribution of examples drawn D_{adv} and D_{train} using a consistent statistical test T .

We validate this hypothesis in Section 5. We show that as few as **50 misclassified adversarial examples** per class are sufficient to observe a measurable difference between legitimate trainings points and adversarial examples for existing adversarial example crafting algorithms.

3.3 Detecting Adversarial Examples

The main limitation of statistical tests is that they cannot detect adversarial examples on a per-input basis. Thus, the defender must be able to **collect a sufficiently large batch of adversarial inputs before it can detect the presence of adversaries.** The defender can **uncover the existence of malicious behavior (as would an intrusion detection system) but cannot identify specific inputs that were manipulated by the adversary among batches of examples sampled (the specific intrusion).** Indeed, sampling a single example will not allow us to confidently estimate its distribution with a statistical test.

This may not be acceptable in security-critical applications. A statistical test, itself, is therefore not always suitable as a defensive mechanism.

However, we propose another approach to leverage the fact that D_{adv} is different from D_{train} . We augment our learning model with an additional *outlier class* C_{out} . We then train the ML model to classify adversarial examples in that class. Technically, C_{out} thus contains all examples that are not drawn from any of the learned distributions $D_{\text{train}}^{C_i}$. We seek to show that this augmented classifier can detect newly crafted adversarial examples at test time.

HYPOTHESIS 2. The augmented classifier with an outlier class C_{out} successfully detects adversarial examples

We validate Hypothesis 2 in Section 6. We also address the potential existence of an arms race between attackers and defenders in Section 7.

4 EXPERIMENTAL SETUP

We describe here the experimental setup used in Sections 5, 6 and 7 to validate the hypotheses stated in Section 3. Specifically, we design our setup to answer the following experimental questions:

- **Q1: How well do statistical tests distinguish adversarial distributions from legitimate ones?** In Section 5, we first find that the MMD and energy distance can statistically distinguish adversarial examples from legitimate inputs. Statistical tests can thus be designed based on these metrics to detect adversarial examples crafted with several known techniques. In fact, we find that often a sample size of 50 is enough to identify them.
- **Q2: Can detection be integrated in ML models to identify individual adversarial examples?** In Section 6, we show that classifiers trained with an additional outlier class detects > 80% of the adversarial examples it is presented with. We also find that in cases where malicious inputs are undetected, the perturbation introduced to evade the model needs to be increased by 150%, making the attack more expensive for attackers.
- **Q3: Do our defenses create an arms race?** We also find that our model with an outlier class is robust to adaptive adversaries, such as the ones using black-box attacks. Even when such adversaries are capable of closely mimicking our model to perform a black-box attack, they are still detected with 60% accuracy in the worst case, and in many cases with accuracies larger than > 90%.

To answer these questions comprehensively, we use several datasets, models and adversarial example algorithms in an effort to represent the ML space. We will introduce them in more detail in the next section.

4.1 Adversarial example crafting

In our experiments, we consider the following attacks. Before we describe them, we want to remark that we do not consider functionality or utility of these attacks, in an attempt to study a worst case scenario.

Fast Gradient Sign Method (FGSM)—This attack computes the gradient of the model’s output with respect to its input. It then perturbs examples in that direction. The computational efficiency of this attack comes at the expense of it introducing large perturbations that affect the entire input. This attack is not targeted towards a particular class. We used the initial implementation provided in the `cleverhans` v.0.1 library [12] and varied the perturbation in the experiments.

Jacobian-based Saliency Map Approach (JSMA)—In contrast to the FGSM, this attack iteratively computes the best feature to perturb for misclassification as a particular (usually closest) class. This yields an adversarial example with fewer modified features, at the expense however of a higher computational cost. We rely again on the implementation provided in the `cleverhans` v.1 library [12].

SVM attack— This attack is described in [26]. It targets a linear SVM by shifting the point orthogonally along the decision boundary. The result is a perturbation similar to the one found by the FGSM. In the case of SVMs however, the perturbation depends on the target class.

Decision Tree (DT) attack— We implemented a variant of the attack from [26] where we search the shortest path between the leaf in which the sample is currently at and the closest leaf of another class. We then perturb the feature that is used in the first common node shared by the two paths. By repeating this process, we achieve a misclassification. This attack modifies only few features and is not targeted.

4.2 Datasets

We evaluate our hypothesis on three datasets.

MNIST—This dataset consists of black-and-white images from 0 to 9 taking real values [19]. It is composed of 60,000 images, of which 10,000 form a test dataset. Each image has 28x28 pixels.

DREBIN—This malware dataset contains 545,333 binary malware features [2]. To make adversarial example crafting faster, we apply dimensionality reduction, as done by Grosse et al. [14], to obtain 955 features. The dataset contains 129,013 Android applications, of which 123,453 are benign and 5,560 are malicious. We split this dataset randomly in training and test data, where the test data contains one tenth of all samples.

Due to its binary nature, it is straightforward to detect attacks like the FGSM or the SVM attack: they lead to non-binary features. We did, nonetheless, compute them in several settings to investigate performance of the detection capabilities. We did not restrict the features that can be perturbed (in contrast to previous work [14]) in an effort to evaluate against stronger adversaries.

MicroRNA—This medical dataset consists of 3966 samples, of which 1280 are breast cancer serums and the remaining are non-cancer control serums. We restrict the features to the 5 features reported as most useful by the original authors [1]. When needed, we split the dataset randomly in training and test subsets, with a 1/10 ratio. This dataset contains real-valued features, each with different mean and variance. We computed perturbations (for SVM, the FGSM and the JSMA) dependent on the variance of the feature to be perturbed.

4.3 Models

We now describe the details of the models used. These models were already introduced in Section 2.1, and their implementation available at [URL](https://github.com/0x09b) blinded.

Decision Trees—We use the Gini impurity as the information gain metric to evaluate the split criterion.

Support Vector Machines—We use a linear multi-class SVM. We train it with an l2 penalty and the squared hinge loss. When there are more than two classes, we follow the one-vs-rest strategy.

Neural Networks—For MNIST, the model has two convolutional layers, with filters of size 5x5, each followed by max pooling. A fully connected layer with 1024 neurons follows.

The DREBIN model reproduces the one described in [14]. It has two fully connected layers with 200 neurons each. The network on the MicroRNA data has a single hidden layer with 4 neurons.

All activation functions are ReLU. All models are further trained using early stopping and dropout, two common techniques to regularize the ML model’s parameters and thus improve its generalization capabilities when the model makes predictions on test data.

Logistic regression—We train a logistic regression on MicroRNA data with dropout and a cross-entropy loss.

Most of these classifiers achieve accuracy comparable to the state-of-the-art on MNIST⁴. On DREBIN, the accuracy is larger than 97.5%. On MicroRNA, the neural network and logistic regression achieve 95% accuracy.

5 IDENTIFYING ADVERSARIAL EXAMPLES USING STATISTICAL METRICS AND TESTS

We answer the first question from Section 4: *in practice, how well do statistical tests distinguish adversarial distributions from legitimate ones?* We find that two statistical metrics, the MMD and the energy distance, both reflect changes—that adversarial examples make to the underlying statistical properties of the distribution—by often strong variations of their value. Armed with these metrics, we apply a statistical test. It detects adversarial examples confidently, even when presented with small sample sets. This validates Hypothesis 1 from Section 3: adversarial examples exhibit statistical properties significantly different from legitimate data.

5.1 Characterizing adversarial examples with statistical metrics

We consider two statistical distance measures commonly used to compare higher dimensional data: (1) the maximum mean discrepancy, and (2) the energy distance.

Maximum Mean Discrepancy (MMD)—Recall from Section 2.3 that this divergence measure is defined as:

$$\text{MMD}_b[\mathcal{F}, X_1, X_2] = \sup_{f \in \mathcal{F}} \left(\frac{1}{n} \sum_{i=1}^n f(x_{1i}) - \frac{1}{m} \sum_{i=1}^m f(x_{2i}) \right)$$

where $x_{1i} \in X_1$ is the i -th data point in the first sample. $x_{2j} \in X_2$ is the j -th data point in the second sample, which is possibly drawn from another distribution than X_1 . $f \in \mathcal{F}$ is a kernel function chosen to maximize the distances between the samples from the two distributions. In our case, a Gaussian kernel is used.

Energy distance (ED)—The ED, which is also used to compare the statistical distance between two distributions, was first introduced by Székely et al. [35]. It is a specific case of the maximum mean discrepancy, where one does not apply any kernel.

Measurement results—We perturb the training distribution of several MNIST models (a neural network, decision tree, and support vector machine) using the adversarial example crafting algorithm presented in Section 4.1 that is suitable for each model. We then measure the statistical divergence (i.e. the distance) between the adversarially manipulated training data and the model’s training data by computing the MMD and ED.

All data points are drawn randomly out of the 60,000 training points. The chance of having a particular sample and its modified

⁴The linear SVM and decision tree only achieve 92.7% and 67.4% on MNIST.

Manipulation	Parameters	MMD	ED
<i>Original</i>	-	0.105	130.85
FGSM	$\epsilon = 0.07$	0.281	157.904
FGSM	$\epsilon = 0.275$	0.603	213.967
JSMA	-	0.14	137.63
DT attack	-	0.1	130.71
SVM attack	$\epsilon = 0.25$	0.524	186.32
Flipped	-	0.306	135.0
Subsampling	45 pixel	2.159	102.7
Gaussian Blur	4 pixel	1.021	128.52

Table 1: Maximum mean discrepancy (MMD) and energy distance (ED) between the original distribution and transformed distributions obtained by several adversarial and geometric techniques on MNIST. Values are averaged over sets of 1,000 inputs sampled randomly from the particular data. For each technique, parameters such as the perturbation magnitude for the FGSM or the number of blurred pixels are given. The JSMA leads to a change of on average 20 pixels, whereas the DT attack changes on average 1 pixel.

counterpart in the same batch is thus very small. To give a baseline, we also provide the distance between the unmodified training and test distributions. At this point, we do not provide the variances, since we consider this to be a sanity check for the following steps. We present the results of our experiments in Table 1.

We observe that for most adversarial examples, there is a strong increase in values of the MMD and ED. In the case of the FGSM, we observe that the increase is stronger with larger perturbations ϵ . For the JSMA and the DT attack, changes are more subtle because these approaches only modify very few features.

We then manipulate the test data using geometric perturbations. While these are not adversarial, they are nevertheless helpful to interpret the magnitude of the statistical divergences. Perturbations considered consist in mirroring the sample, subsampling from the original values, and introducing Gaussian blur.⁵ We find that mirroring and subsampling affect both the MMD and ED, whereas Gaussian blur only significantly increases the ED.

In this first experiment, we observed that there exists measurable statistical distances between samples of benign and malicious inputs. This justifies the design of consistent statistical tests to detect adversarial distributions from legitimate ones.

5.2 Detecting adversarial examples using hypothesis testing

We apply a statistical test to evaluate the following hypothesis: *samples from the test distribution are statistically close to samples from the training distribution*. We expect this hypothesis to be accepted for samples from the legitimate test distribution, but rejected for samples containing adversarial examples. Indeed, we observed in Section 5.1 that adversarial distributions statistically diverge from the training distribution.

⁵This geometric perturbation approximates an attack against ML models introduced by Biggio et al. [4]. Indeed, the adversarial inputs produced by this attack appear as blurry, with less crisp shapes.

Two-sample hypothesis testing— As stated before, the test we chose is appropriate to handle high dimensional inputs and small sample sizes.⁶ We compute the biased estimate of MMD using a Gaussian kernel, and then apply 10,000 bootstrapping iterations to estimate the distributions. Based on this, we compute the p-value and compare it to the threshold, in our experiments 0.05. For samples of legitimate data, the observed p-value should always be very high, whereas for sample sets containing adversarial examples, we expect it to be low—since they are sampled from a different distribution and thus the hypothesis should be rejected.

The test is more likely to detect a difference in two distributions when it considers samples of large size (i.e., the sample contains more inputs from the distribution).

Whenever we write *confidently detected* at sample size x , we mean that all 200 instances of the test on x randomly sampled examples from each of the two distributions rejected H_0 . Percentages reported correspond to the n times of 200 the test accepted H_0 .

Results— Regardless of the sample size, the hypothesis acceptance for benign data generally lies around 95%. This means that the benign data is confidently identified as such. The sample size (i.e., the number of adversarial examples) required to confidently detect adversarial distributions is given in Table 1a for the three datasets (MNIST, DREBIN, MicroRNA).

For most datasets and models, a sample size of 50 adversarial examples is sufficient for the statistical test to reject H_0 when comparing this sample to a sample from the benign distribution. Thus, the statistical test identifies adversarial examples with strong confidence, despite having few points of comparison (relatively to the training set size of 50,000 for MNIST).

Some exceptions should be noted. A sample of at least 100 inputs is required to confidently detect adversarial examples crafted with the JSMA on a neural network. In addition, the test is unable to detect adversarial examples crafted for MNIST on decision trees.

However, these two observations are consistent with results from Section 5.1, which showed that these attacks yielded adversarial inputs with less distinguishable statistical properties than the FGSM for instance. Another result confirms this explanation: the SVM attack, which was observed to lead to large changes in the MMD and energy distance, is as well easily detected by the two-sample test. A sample of 10 adversarial inputs is sufficient to confidently reject the H_0 .

Briefly put, these results support that the distribution of JSMA or FSGM adversarial examples differs from legitimate inputs, where we used a statistical hypothesis test on a set of inputs.

5.3 Decreasing sample size by performing class-wise statistical tests

In an effort to reduce the number of inputs (i.e., the sample size) required for the statistical test to be confident, we propose here to refine it. We previously assumed that the dataset was generated by a single underlying distribution. However, each class of the problem can be viewed as its own distribution. We thus perform below the statistical test on samples grouped according to their label. We separate the inputs according to their original label, or the (wrong)

⁶We used an implementation publicly available at https://github.com/emanuele/kernel_two_sample_test.

Dataset	FGSM	JSMA	SVM	DT
MNIST	50 (.275)	100 (16)	10 (.25)	- (1)
DREBIN	50 (.6)	50 (2)	10 (.25)	50 (2)
Micro	50 (.6)	10 (3)	*	50 (1)

(a) Whole datasets. The average adversarial perturbation introduced is characterized in parenthesis either by stating the perturbation parameter ϵ (FGSM, SVM attack), or the number of perturbed features (JSMA, DT attack).

Attack:	FGSM		JSMA		SVM		DT	
Class:	O	P	O	P	O	P	O	P
MNIST	50	50	50	100	10	10	50	-
DREBIN (+)	10	10	50	50	10	10	50	-
DREBIN (-)	10	10	50	50	10	10	10	50
Micro (+)	10	10	10	10	*	*	10	10
Micro (-)	10	10	10	10	*	*	10	50

(b) The statistical test is run either with the original class (O) of the input or the class predicted by the model to the perturbed input (P). Upper row for DREBIN refers to malware class (+), second to benign programs (-). For MicroRNA, (+) are the cancer serum, (-) is the control group. for MNIST, we report average values over all classes.

Figure 1: Minimum sample size (i.e., number of adversarial inputs) needed to confidently detect adversarial examples. Stars indicate experiments that were not conducted because the attack failed to succeed (yielding reductions in accuracy smaller than $< 30\%$) or initial accuracy was too low. In other cases (-), even a sample size of 500 was not enough to detect the adversarial examples.

label assigned by the model. These tests are found to be confident for smaller sample sizes.

Results— Experimental results are given in Table 1b. On MNIST, we find that these new class-wise statistical tests reduce the sample size needed to detect the JSMA to 50 examples. This is also the case with DREBIN, where the sample size is reduced to 10. For the MicroRNA we only observe a change in sample size concerning the decision tree attack, where again confident detection is already possible at a sample size of 10.

For all datasets, we find that using a statistical test based on the distribution of the class in which the inputs are (wrongly) misclassified is more effective than using the class from which they were derived: the latter even completely fails for decision trees on MNIST. These results are consistent across the two other datasets (DREBIN and MicroRNA). In Figure 2, however, we observe that on MNIST when testing for the FGSM examples, this tendency is reversed.

Briefly put, we observed that the minimum sample size to achieve confident detection with a class-wise test is smaller than the sample size required by the general statistical test. We also noted that statistical tests comparing adversarial examples with training examples

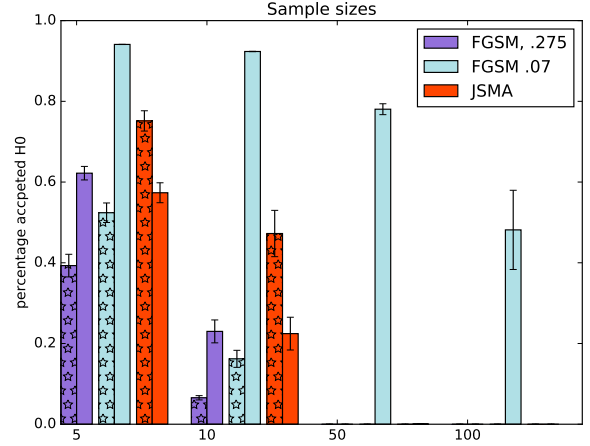


Figure 2: Frequency of hypothesis H_0 acceptance with respect to the sample size (number of adversarial examples) on which the statistical test is performed. Lower values indicate that the hypothesis was rejected in more cases (e.g., the adversarial examples are detected as outside the expected distribution), which is the defender’s goal. The statistical test is defined either with the original class of the input (bars with patterns), or the class predicted by the model on the perturbed input (plain bars).

from the class they are misclassified as (rather than the class they were derived from) are more confident.

To close this section, we want to remark that a possible conclusion from the findings in this section is to apply statistical outlier detection to detect adversarial examples. This yields a model agnostic way to detect adversarial examples. We experimented with simple outlier detection models and found, however, that many of them were not able to handle the high dimensional data with good confidence⁷. Since the classifiers themselves however can also be trained to perform outlier detection, we went for the approach described in the following section.

6 INTEGRATING OUTLIER DETECTION IN MODELS

In the previous section, we concluded that the distribution of adversarial examples statistically differs from the expected distribution. Yet, the confidence of the test diminishes with the number of examples in the sample set analyzed: this test cannot be used to identify which specific inputs are adversarial among a set of inputs.

In this Section, we provide an answer to our second experimental question: “Can we detect individual adversarial examples?” Our approach adds an additional output to the model. The model is trained to assign this new output class to all adversarial inputs. In other words, we explicitly train models to label all inputs that are not part of the expected distribution as part of a new *outlier class*.

⁷We used the Two-Sample-Kernel Test with a single sample and Tukey’s test. We further investigated several threshold/quartile based combinations for a radial, linear, and Gaussian distances and kernels.

In the following experiments, we show that this approach is complementary to the statistical test introduced in Section 2.3 because it enables the defender to accurately identify whether a given input is adversarial or not.

Intuition— In the previous section, we have shown that the feature distribution of adversarial examples differs significantly from the distribution of benign training data. Yet, there exists no real feature distribution D_{real} for adversarial examples: they are instead derived from the feature distribution of the original classes through minimal perturbation based on reconnaissance of the attacked classifier’s behavior.

In the following, we want to leverage this insight while the classifier is being trained. Our goal is to be able to detect individual adversarial examples, as discussed in Section 3. Since the distribution drift between the training and adversarial test distributions is detectable, we can hypothesize that it is learnable as well. Assuming that the classifier generalizes well to adversarial examples it has not seen during training, this would enable us to detect adversarial examples.

Training with an outlier class— We start the process by training an initial model N_{in} on the original data $D = \{X, Y\}$. We compute adversarial examples for N_{in} on the training data, denoted as X_{in} . We then train a new model, N_{p1} on an augmented dataset, $X \cup X_{in}$, where all adversarial examples are assigned to the outlier class. In particular, adversarial examples of different crafting algorithms are in the same class. Specifically, we arrange batches of inputs analyzed by the learning algorithm such that 2/3 are legitimate training examples from X and 1/3 are adversarial from X_{in} .

To evaluate the robustness of this model, we craft adversarial examples on the test data $D_t = \{X_t, Y_t\}$ for the augmented model. In the case of JSMA, the target class will always be selected as the second most confident class, except this class is the outlier class. We thus always craft towards the strongest original class. We then report the misclassification or error on adversarial examples, e.g. all adversarial examples that were not classified as outliers or as their original class. Additionally, we report the percentage of all adversarial examples that were classified as outliers.

Performance impact— In order for the outlier class to be a practical mechanism for detecting adversarial inputs, the impact on the model accuracy on benign test data should be limited.

On MNIST, we observe that the parameters of our model augmented with an outlier class converge slightly slower (in terms of training iterations required) to their final values compared to an unmodified model. In addition, the model’s accuracy varies between 98.0% and 98.8% when it is presented with legitimate inputs. Thus deploying the defense comes at the cost of a decrease from the original model accuracy of 99.2%.

On DREBIN, the original network achieves an accuracy on the benign test data of 97.7%. Considering the mean, we observe a slight decrease of 0.2% in accuracy for the model trained using an additional class.

On the microRNA data, the original logistic regression achieves an accuracy around 95.4%. With an outlier class, accuracy drops to 93.4%. The same holds for neural networks trained on the same dataset.

Attack	ϵ	Recovered	Detected	Error
FGSM	0.2	0.2%	99.78%	0.02%
FGSM	0.28	0.01%	99.91%	0.08%
FGSM	0.3	0.01%	99.93%	0.06%
FGSM	0.4	0%	99.99%	0.01%
FGSM	> 0.5	0%	100%	0%
JSMA	≤ 80	0.02%	83.76%	16.22%

Table 2: Adversarial example detection rate for models trained with an outlier class on MNIST. All models are trained on $\epsilon = 0.275$ for the FGSM and $\epsilon \leq 200$ for the JSMA. The first columns indicate the attack used upon completion of training and its parameter. *Recovered* indicates the rate of adversarial examples classified in the original class of the input they were crafted from. *Detected* indicates the percentage of adversarial examples that were classified as outliers. The *error* rate is simply the remaining adversarial examples (those not correctly classified or detected as outliers).

Detection results— In this series of experiments, we first measure the detection performance of the model on malicious examples. We then study the failure cases of our mechanism on benign data.

Results on MNIST is given in Table 2. A large proportion of adversarial examples are detected or their original class is recovered (so they are correctly classified). The outlier class mechanism can be successfully trained to detect adversarial examples produced by both the FGSM and JSMA attacks, and we discuss later in Section 7 how it fares with mixtures of both attacks. Generally speaking, the detection rate increases with the perturbation magnitude, while the recovery rate decreases.

We now report the results on the DREBIN dataset. Concerning the FGSM adversarial examples, we observed, independently from chosen ϵ , an misclassification around 92.3%. The network could further not be hardened against those adversarial examples: after training, the accuracy was still 92.2%. For the JSMA, we observe an initial misclassification of 99.991% by changing 2.3 features.⁸ When retraining on adversarial examples, we do not observe any increase in robustness. We do observe, however, an increase in the number of changed features up to 5.8 when trained on JSMA examples.

To understand whether the limited effectiveness of our defense on DREBIN is a consequence of the binary nature of its data or the stronger success of the attack, we implemented a second attack with a worse heuristic that initially modifies 6.6 features on average. By training on adversarial examples crafted with this defense, it became impossible to craft adversarial examples using the same modified JSMA with an upper limit of 90 changed features.

We also trained a simple logistic regression on the MicroRNA data and trained it on adversarial examples. The results are depicted in Table 3. We only applied the FGSM, since the JSMA was not successful (only 40% of the adversarial examples evaded the model). Initially, misclassification was 95.7% by applying a perturbation of $\epsilon = 1.0$. Further increase of ϵ had no effect. Training logistic regression with an outlier class on adversarial examples given $\epsilon = 1.0$, we

⁸This initial perturbation is much less than reported in the original work, since we do not restrict the features as done previously [14].

ϵ	log reg	log reg+1		NN+1
	Accuracy	Error	Detected	Error
0.2	87.4%	3.5%	16.7%	4.5%
0.4	42.9%	5.9%	64.6%	10.4%
0.6	15.9%	32.3%	64.9%	4.5%
0.8	5.1%	69.2%	26.3%	2%
1.0	4.3%	87.2%	6.7%	0.7%

Table 3: Accuracy and detection rate of MicroRNA logistic regressions (log reg) and neural networks (NN). We present a baseline model and two models trained with additional class (+1), both trained on FGSM at $\epsilon = 1.0$ and $\epsilon = 0.8$. Attack parameter for FGSM is given in the first column. Error refers to the percentage of misclassified adversarial examples.

obtain a misclassification of 87.2%. For lower perturbations, we can decrease misclassification, however. The limited improvement is most likely due to the limited capacity of logistic regression models, which prevents them from learning models robust to adversarial examples [11]. Thus, we trained a neural network on the data. Initially it could be evaded with the same perturbation magnitude and success. Yet, when trained with the outlier class, misclassification was 0.7% on the strongest ϵ .

Wrongly classified benign test data—Next, we investigate the error cases of our mechanism. The number of false positives, benign test examples of the original data that are wrongly classified as outliers, represents a small percentage of inputs: e.g., 0.5% on MNIST. In addition, we draw confusion matrices for the benign test data in Figure 3. The diagonal indicates correctly classified examples and is canceled out to better visualize out-of-diagonal and misclassified inputs. Misclassification between classes is very similar in the original case and when training with the FGSM. In the interest of space, we thus omit the confusion matrix for original data. In contrast, when training on JSMA examples, a large fraction of misclassified data points is no longer misclassified as a legitimate class, but wrongly classified as outliers.

7 PREVENTING THE ARMS RACE

A key challenge in ML security lies in the fact that no defense guarantees resilience to future attack designs. This contrasts with ML privacy where differential privacy guarantees withstand all hypothetical adversaries. Such an arms race may only be broken by mechanisms that have been proven to be secure in an expressive security model, such as the one of differential privacy.

While providing any formal guarantees for the methods proposed here is intrinsically hard given the nature of optimization problems solved by ML algorithms, we evaluate here their resilience to adaptive strategies. We first show that the statistical test still performs well when presented with a mixture of benign and malicious inputs. We then demonstrate the robustness of our models augmented with the outlier class to powerful black-box strategies that have evaded previous defenses.




Figure 3: Confusion matrices on benign test inputs of MNIST. The horizontal axis denotes the original label, and the vertical one the output class of the network. Left side corresponds to training on FGSM examples, right side on JSMA examples. The diagonal (correctly classified data points) has been zeroed, indices correspond to MNIST classes, where 10 is the outlier class. Both matrices are normalized in the same scale, i.e., same color means same number of misclassified samples. Brighter indicates a higher number of misclassified examples.

7.1 Robustness of the Statistical Test

Though we introduced the statistical test as a defense, but as a tool to investigate the distribution of adversarial examples, one might perform it on a batch before submitting inputs to the ML model. A natural question is then whether an adversary aware of this defense could evade it by constructing adversarial examples simultaneously misleading the model and the statistical test.

Theoretically, a statistical test cannot be misled because adversarial examples necessarily deviate from the expected distribution (see Section 3). Yet, this assumes that the defender is capable of running statistical tests on sufficiently large sample sets of inputs. In that case, it is guaranteed that the null hypothesis would be rejected. However, this may not always be the case in practice as such sampling may require a potentially infinite number of inputs. As such, we measure the confidence of our statistical test as it is presented with more realistic sample sets of inputs in an effort to demonstrate its robustness.

We consider two scenarios where the attacker adapted its strategy. First, the attacker might hide small numbers of adversarial examples among a large number of benign samples. In the following, we thus investigate the statistical test’s performance in detecting adversarial inputs in the presence of legitimate inputs. Second, we consider sample sets of adversarial examples, where the adversary executed more than one adversarial crafting algorithm, like for instance the FGSM and the JSMA.

Mixture of adversarial examples—We observe reductions in the detection confidence for mixtures of adversarial examples. This is in particular the case when one of the adversarial example kind is hard to detect. For instance, we observed in Section 5.2 that adversarial inputs for decision trees are hard to detect on MNIST. This reduces the performance of the statistical test on samples that contain these examples. Our full results are depicted in Figure 4.




Figure 4: Evaluating mixtures of different adversarial examples and benign data. X-axis indicates the percentage of benign data. Y-axis is the percentage where the statistical test confirms that the data is from the same distribution (lower is better). We observe that the more benign examples (moving to the right), the harder it is to detect the remaining adversarial examples.

Mixtures of adversarial and legitimate inputs—The trade-off between the ratio of benign and adversarial examples, test confidence and sample size is shown in Figure 4. The test is more confident when the percentage of adversarial examples is high or the sample size is large. Hence, one is less likely to detect adversarial examples mixed with legitimate inputs among small sets of inputs.

We conclude that the statistical test’s confidence decreases, as expected, when the adversary submits very few adversarial examples among large sets of legitimate inputs. This is however the main motivation behind the outlier class mechanism introduced in Section 6.

7.2 Robustness of the Outlier Class

We now investigate the performance of models augmented with an outlier class in the face of adversaries aware of that defense. We first show that these models are able to generalize to varying attacker strategies, i.e., they can detect adversarial inputs crafted using a different algorithm than the one used to train the outlier class.

In addition, we note that in all previous experiments, we considered adversaries directly computing adversarial examples based on the defended model’s parameters. Instead, we here evaluate the model’s robustness when attacked using black-box strategies. These powerful attacks have been shown to evade previously proposed defenses, such as adversarial training and defensive distillation [25]. The reason is that these defenses did not actually fix model errors but rather manipulated the model’s gradients, thus only making it harder for the adversary to craft adversarial examples when they are computed directly on the targeted model.

Training		Attack			
ϵ	Attack	ϵ	R	D	Error
≤ 200	JSMA	0.1	2.04%	77.16%	20.8%
≤ 200	JSMA	0.275	2.07%	96.6%	2.95%
≤ 200	JSMA	0.4	0.22%	98.45%	1.33%
≤ 200	JSMA	0.6	0.13%	99.58%	0.29%
0.275	FGSM	≤ 80	0%	9.63%	90.37%

Table 4: Misclassification and adaptive detection rate DNN trained with an outlier class on MNIST. All models are trained on $\epsilon = .275$ FGSM examples and $\epsilon < 200$ JSMA examples. The attacks used to evaluate the detection performance are different from the one used to train the outlier class. *Recovered* (R) indicates the rate of adversarial examples classified in the original class of the input they were crafted from. *Detected* (D) indicates the percentage of adversarial examples that were classified as outliers. The *error* rate is simply the remaining adversarial examples (those not correctly classified or detected as outliers)

A simple but highly successful strategy is then for the adversary to train an auxiliary model that mimics the defended model’s predictions, and then use the auxiliary model to find adversarial examples that are also misclassified by the defended model. In the following, we show that our models with an outlier class are also robust to such strategies.

Robustness of detection to adaptive attackers— We investigate whether the outlier class generalizes to other adversarial example crafting techniques. In other words, we ask whether defending against one type of adversarial examples is sufficient to mitigate an *adaptive* attacker using multiple techniques to craft adversarial examples.

We thus trained the MNIST model’s outlier class with only one kind of adversarial example, and then observed its robustness to another kind of adversarial examples. Table 4 reports this result for varying attack parameter intensities. If we train the model on JSMA adversarial examples, it is also robust to adversarial FGSM examples crafted. If perturbations are high, misclassification is smaller than 3% percent. The reverse case, however, does not hold: a model trained on FGSM is only slightly more robust than the original model.

We did not perform this experiment on DREBIN or MicroRNA datasets, since we could only apply one attack on each of them.

Robustness of detection to black-box attacks performed using transferability— We now show that our proposed outlier class mechanism is robust to an additional attack vector against ML models: black-box attacks exploiting adversarial example transferability. These techniques allow an adversary to force a ML model to misclassify without knowledge of its model parameters (and sometimes even without knowledge of its training data) by computing adversarial examples on a different model than the one targeted [11, 25, 34].

In order to simulate the *worst-case* adversary, we train the *substitute* model from which we will transfer adversarial examples back to

ϵ	Attack	R	D	E
0.1	FGSM	42.67%	55.72%	1.61%
0.275	FGSM	0%	100%	0%
> 0.4	FGSM	0%	100%	0%
≤ 80	JSMA	0.3%	97%	2.7%

ϵ	Attack	R	D	E
0.1	FGSM	17.64%	81.64%	0.72%
0.275	FGSM	0%	99.98%	0.02%
> 0.4	FGSM	0%	100%	0%
≤ 80	JSMA	0.46%	92.85%	6.69%

Table 5: Robustness of a MNIST model (with an outlier class) to black-box attackers. All models are trained on $\epsilon = .275$ FGSM examples and $\epsilon < 200$ JSMA examples. *Recovered (R)* indicates the rate of adversarial examples classified in the original class of the input they were crafted from. *Detected (D)* indicates the percentage of adversarial examples that were classified as outliers. The *Error (E)* rate is simply the remaining adversarial examples (those not correctly classified or detected as outliers). Above Table show an attacker using a substitute model without an outlier class. Lower one indicates an attacker using a substitute model with an outlier class, equivalent to tested model.

the victim model with the same training data than the one used by the victim. Again, a black-box attack may be possible even without such knowledge [25]. Yet, this allows us to consider a particularly strong adversary capable to closely mimic our model. We further focus on the robustness of our approach to such attacks, rather than a demonstration of the feasibility of such attacks. We train two substitute models, one including the outlier class (denoted BB+1) as does the victim, the other not (BB). We then compute adversarial examples on the substitute model and evaluate the misclassification rates of the victim model.

MNIST results— The results on MNIST are given in Table 5. We observe in both settings high robustness against the adversarial examples computed using the FGSM. Depending on the setting, we misclassify not a single of the examples when $\epsilon > 0.4$ for BB+1 and $\epsilon > 0.275$ for BB. We further observe high robustness as well concerning the adversarial examples computed using JSMA: In the setting of BB+1, the misclassification is around 6.7%, for BB even 2.7%. In all cases, most of the adversarial examples are classified as outliers.

DREBIN results— We observe that, though the network is vulnerable to direct attacks, it is much more resistant to adversarial examples crafted on another network. In this setting, we only compute JSMA examples. Given a substitute trained with additional class, we observe misclassification rates of 35.8%, 47.1% are detected as outliers. For JSMA examples crafted on the original network 0.8% are misclassified; 98.6% of adversarial examples are classified as outliers.

MicroRNA results— The results on the MicroRNA data are depicted in Table 6. We observe that logistic regression works well

in detecting smaller perturbations till 0.4. For larger ϵ , however, the misclassification drastically increases, partially to 88%. Again, we assign the small improvement to the limited capacity of logistic regression, we trained additionally neural networks. We observed lower misclassification in settings where the perturbation is maximal ($\epsilon = 1.0$). In general, we observe that logistic regression is more robust in the BB setting, whereas neural networks are more robust in the BB+1 setting (both with some exceptions).

8 DISCUSSION

We discuss the limitations of the mechanisms proposed to detect adversarial examples: statistical testing, and an outlier class. We also explore avenues for future work.

Statistical Test— As we have seen, one of the major strengths of kernel-based statistical tests is that they operate and thus detect the presence of adversarial examples already in feature space, before these inputs are even fed to the ML model. Intuitively, we observed that the larger the perturbation applied is, the more likely it is to be confidently detected by the statistical test. Adversarial example crafting techniques that modify few features (like the JSMA or the decision tree attack) or perturb the features only slightly (small values of ϵ for the FGSM) are less likely to be detected.

This finding is consistent with the underlying stationary assumption made by all ML approaches. Since adversarial examples are not drawn from the same distribution than benign data, the classifier is incapable of classifying them correctly. This property also holds for the training data itself, and as such, we expect it to generalize to poisoning attacks. In such attacks, the adversary attempts to degrade learning by inserting malicious points in the model’s training data. This is however outside the scope of this work, and we leave this question to future work.

Integrating Outlier Detection— We further observe that adding an outlier class to the model yields robustness to adaptive attack strategies, and needed perturbation is increased. Concerning JSMA, for some datasets, we do not achieve robustness to adversarial examples. This most likely depends on the initial vulnerability of the data: For DREBIN we changed barely more than one feature, for MNIST almost twenty. At the same time, MNIST has slightly less features, of which the pixels at the borders are barely used. Thus, having less features and a higher perturbation to learn from might yield larger robustness to adversarial examples. Additionally, further factors might include inter-class and intra-class distances, or the variability of the computed adversarial examples.

Further, the confusion matrices from Figure 3 seem to suggest that FGSM lie in a different halfspace than the original data: the outlier class trained on JSMA examples contains benign data points whereas the FGSM one does not. This might indicate that JSMA examples lie rather between benign classes.

We further observed that knowledge about the attack is not necessarily needed: training on adversarial examples computed using the JSMA hardens against computing FGSM examples. Perhaps surprisingly, this does not hold the other way around. In general, since FGSM is non-targeted and less optimal than JSMA, further work is needed whether the outlier class generalizes from targeted to non-targeted attacks or from more optimal to less optimal attacks. In theory, we could feed samples from the whole feature space

ϵ	BB						BB+1					
	logistic regression			neural network			logistic regression			neural network		
	R	D	E	R	D	E	R	D	E	R	D	E
0.2	84.6%	11.9%	3.5%	83.6%	12.9%	3.5%	81.3%	14.9%	3.8%	80.8%	15.7%	3.5%
0.4	49.7%	46.7%	3.6%	47.7%	48.7%	3.6%	35.1%	59.8%	5.1%	27.8%	65.2%	5%
0.6	21.5%	74.7%	3.8%	15.4%	80.6%	4%	9%	74.5%	16.5%	4.5%	68.9%	26.6%
0.8	14.8%	75.3%	9.8%	7.3%	81.8%	10.9%	6%	55.8%	38.2%	4%	30.8%	65.2%
1.0	9.3%	68.2%	22.5%	3.8%	65.4%	30.8%	4.3%	38.6%	57.1%	6%	6.0%	88%
0.2	83.3%	13.9%	2.8%	82.6%	14.6%	2.8%	76.3%	21.5%	2.2%	84.1%	12.1%	3.8%
0.4	45.2%	54.0%	1.8%	42.9%	55.3%	1.8%	30.3%	68.4%	1.3%	50.3%	45.7%	4%
0.6	23.4%	76.0%	10.6%	16.5%	78.3%	5%	4%	93.2%	2.8%	14.6%	81.6%	3.8%
0.8	13.4%	76%	10.6%	6%	80.3%	13.7%	2%	94.2%	3.8%	3.2%	95.0%	1.8%
1.0	3%	79.8%	17.2%	2%	73.2%	24.8%	2%	94.7%	3.3%	2%	97.2%	0.8%

Table 6: Black box setting for logistic regression (upper part) and a neural network (lower part) on the MicroRNA data trained using the outlier class. Substitutes are logistic regression or a neural network (NN), either trained without (BB) or with an additional class (BB+1). If an outlier class is used, FGSM examples at $\epsilon = 1.0$ and $\epsilon = 0.8$ are used for training. We report parameters of attack (ϵ). *Recovered* (R) indicates the rate of adversarial examples classified in the original class of the input they were crafted from. *Detected* (D) indicates the percentage of adversarial examples that were classified as outliers. The *Error*(E) rate is simply the remaining adversarial examples (those not correctly classified or detected as outliers).

except the location of the classes, and thus obtain a robust classifier without any assumption on the adversary. This is practically infeasible, however. Future work will investigate trade-offs here.

Finally, we want to remark that the benign data, that is labeled as outlier by the network might be beneficial when investigated by an expert[23]. This data might be either excluded from training, or relabeled. This question will be answered in future work.

9 RELATED WORK

Other approaches to detect malicious data points by statistical means have been proposed. However, they all depend on some of the internal activations of deep neural networks models [20, 30, 32]. Hence, these approaches only apply to the specific classifier studied. In contrast, we apply our statistical test directly in feature space, allowing us to propose a model-agnostic detection.

Wang et al. [36] present a similar formal intuition as we do (using an oracle instead of the underlying distribution, though). From this, they formally derive conditions when a classifier is secured against adversarial examples. Further, they proposed a modified version of adversarial training, originally introduced by Goodfellow et al. [11]. In contrast to both of these approaches, we classify adversarial examples in a separate (and additional) outlier class. We also do not compute adversarial examples throughout training but rather use adversarial examples precomputed on a different model before training.

Nguyen et al [24] have introduced an outlier class before. Also Bendale et al [3] propose open networks, that are not confident in their classification all over the feature space. Both, however, do not evaluate and motivate their approach in adversarial settings.

Metzen et al. [18] augment neural networks with an auxiliary network used to detect malicious samples. This additional network shares some of its parameters with the original one, and thus also depends on the features of the network. Further, our outlier class

mechanism is applicable to any ML models. In addition, their approach is limited in settings where the adversary adapts its strategy. Instead, our experiments systematically explore the space of adversaries (with different adversarial example crafting algorithms, datasets and models). We also present a detailed discussion of possible adaptive strategies, such as powerful black-box attacks known to be hard to defend against [25].

10 CONCLUSION

We empirically validated the hypothesis that adversarial examples can be detected using statistical tests before they are even fed to the ML model as inputs. Thus, their malicious properties are model-agnostic.

Furthermore, we show how to augment ML models with an additional class in which the model is trained to classify all adversarial inputs. This results in robustness to adversaries, even those using attack strategies based on transferability—a class of attacks known to be harder to defend against than gradient-based strategies. In addition, when adversarial examples with small perturbations are not detected as outliers, they are original class is often recovered and the perturbed input correctly classified.

Additionally, we expect that combining our approaches together, as well as with other defenses may prove beneficial. For instance, we expect defensive distillation and the statistical test or outlier class to work well together, as defensive distillation has been found to increase the perturbations that an adversary introduces.

ACKNOWLEDGMENTS

Nicolas Papernot is supported by a Google PhD Fellowship in Security. Research was supported in part by the Army Research Laboratory, under Cooperative Agreement Number W911NF-13-2-0045 (ARL Cyber Security CRA), and the Army Research Office under grant W911NF-13-1-0421. The views and conclusions contained

in this document are those of the authors and should not be interpreted as representing the official policies, either expressed or implied, of the Army Research Laboratory or the U.S. Government. The U.S. Government is authorized to reproduce and distribute reprints for government purposes notwithstanding any copyright notation hereon.

This work was supported by the German Federal Ministry of Education and Research (BMBF) through funding for the Center for IT-Security, Privacy and Accountability (CISPA) (FKZ: 16KIS0753).

REFERENCES

- [1] Kawauchi J Takizawa S et al A. shimomura, Shiino S. 2016. Novel combination of serum microRNA for detecting breast cancer in the early stage. *Cancer Sci Mar* 107(3):326-34 (2016).
- [2] Daniel Arp, Michael Spreitzenbarth, Malte Hubner, Hugo Gascon, and Konrad Rieck. 2014. DREBIN: Effective and Explainable Detection of Android Malware in Your Pocket.. In *Proceedings of the 2014 Network and Distributed System Security Symposium (NDSS)*.
- [3] Abhijit Bendale and Terrance E. Boult. 2016. Towards Open Set Deep Networks. In *2016 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2016, Las Vegas, NV, USA, June 27-30, 2016*. 1563–1572. <https://doi.org/10.1109/CVPR.2016.173>
- [4] Battista Biggio, Blaine Nelson, and Pavel Laskov. 2012. Poisoning attacks against support vector machines. *arXiv preprint arXiv:1206.6389* (2012).
- [5] Olivier Bousquet and André Elisseeff. 2002. Stability and Generalization. *The Journal of Machine Learning Research* 2 (2002), 499–526.
- [6] Michael Brückner and Tobias Scheffer. 2011. Stackelberg games for adversarial prediction problems. In *Proceedings of the 17th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, San Diego, CA, USA, August 21-24, 2011*. 547–555. <https://doi.org/10.1145/2020408.2020495>
- [7] Nicholas Carlini and David Wagner. 2016. Towards Evaluating the Robustness of Neural Networks. *CoRR abs/1608.04644* (2016). <http://arxiv.org/abs/1608.04644>
- [8] N. Carlini and D. Wagner. 2017. Adversarial Examples Are Not Easily Detected: Bypassing Ten Detection Methods. *ArXiv e-prints* (May 2017). [arXiv:cs.LG/1705.07263](http://arxiv.org/abs/1705.07263)
- [9] Nilesh Dalvi, Pedro Domingos, Sumit Sanghai, Deepak Verma, et al. 2004. Adversarial classification. In *Proceedings of the tenth ACM SIGKDD international conference on Knowledge discovery and data mining*. ACM, 99–108.
- [10] J.H. Friedman and L.C. Rafsky. 1979. Multivariate generalizations of the Wald-Wolfowitz and Smirnov two-sample tests. *Ann. Stat.; (United States)* 7:4 (Jan 1979).
- [11] Ian J Goodfellow et al. 2015. Explaining and Harnessing Adversarial Examples. In *Proceedings of the 2015 International Conference on Learning Representations*.
- [12] Ian J. Goodfellow, Nicolas Papernot, and Patrick D. McDaniel. 2016. cleverhans v0.1: an adversarial machine learning library. *CoRR abs/1610.00768* (2016). <http://arxiv.org/abs/1610.00768>
- [13] Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander J. Smola. 2012. A Kernel Two-Sample Test. *Journal of Machine Learning Research* 13 (2012), 723–773. <http://dl.acm.org/citation.cfm?id=2188410>
- [14] Kathrin Grosse, Nicolas Papernot, Praveen Manoharan, Michael Backes, and Patrick McDaniel. 2016. Adversarial Perturbations Against Deep Neural Networks for Malware Classification. *CoRR abs/1606.04435* (2016). <http://arxiv.org/abs/1606.04435>
- [15] Peter Hall and Nader Tajvidi. 2002. Permutation tests for equality of distributions in high-dimensional settings. *Biometrika* 89, 2 (2002), 359–374.
- [16] Harold Hotelling. 1931. The Generalization of Student's Ratio. *Ann. Math. Statist.* 2, 3 (08 1931), 360–378. <https://doi.org/10.1214/aoms/1177732979>
- [17] Ling Huang, Anthony D. Joseph, Blaine Nelson, Benjamin I. P. Rubinstein, and J. D. Tygar. 2011. Adversarial machine learning. In *Proceedings of the 4th ACM Workshop on Security and Artificial Intelligence, AISec 2011, Chicago, IL, USA, October 21, 2011*. 43–58. <https://doi.org/10.1145/2046684.2046692>
- [18] Volker Fischer Jan Hendrick Metzen, Tim Genewein and Bastian Bischoff. 2016. On detecting Adversarial Perturbations. (2016). (to appear).
- [19] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. 1998. Gradient-Based Learning Applied to Document Recognition. *Proc. IEEE* 86, 11 (November 1998), 2278–2324.
- [20] Xin Li and Fuxin Li. 2016. Adversarial Examples Detection in Deep Networks with Convolutional Filter Statistics. *CoRR abs/1612.07767* (2016). <http://arxiv.org/abs/1612.07767>
- [21] Wei Liu and Sanjay Chawla. 2010. Mining adversarial patterns via regularized loss minimization. *Machine Learning* 81, 1 (2010), 69–83. <https://doi.org/10.1007/s10994-010-5199-2>
- [22] Daniel Lowd and Christopher Meek. 2005. Good Word Attacks on Statistical Spam Filters. In *CEAS 2005 - Second Conference on Email and Anti-Spam, July 21-22, 2005, Stanford University, California, USA*. <http://www.ceas.cc/papers-2005/125.pdf>
- [23] Brad Miller, Alex Kantchelian, Michael Carl Tschantz, Sadia Afroz, Rekha Bachwani, Riyaz Faizullahoy, Ling Huang, Vaishal Shankar, Tony Wu, George Yiu, Anthony D. Joseph, and J. D. Tygar. 2016. Reviewer Integration and Performance Measurement for Malware Detection. In *Detection of Intrusions and Malware, and Vulnerability Assessment - 13th International Conference, DIMVA 2016, San Sebastián, Spain, July 7-8, 2016, Proceedings*. 122–141. https://doi.org/10.1007/978-3-319-40667-1_7
- [24] Anh Mai Nguyen, Jason Yosinski, and Jeff Clune. 2015. Deep neural networks are easily fooled: High confidence predictions for unrecognizable images. In *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2015, Boston, MA, USA, June 7-12, 2015*. 427–436. <https://doi.org/10.1109/CVPR.2015.7298640>
- [25] Nicolas Papernot, Patrick McDaniel, Ian Goodfellow, Somesh Jha, and al. 2016. Practical Black-Box Attacks against Deep Learning Systems using Adversarial Examples. *arXiv preprint arXiv:1602.02697* (2016).
- [26] Nicolas Papernot, Patrick McDaniel, and Ian J. Goodfellow. 2016. Transferability in Machine Learning: from Phenomena to Black-Box Attacks using Adversarial Samples. *CoRR abs/1605.07277* (2016). <http://arxiv.org/abs/1605.07277>
- [27] Nicolas Papernot, Patrick McDaniel, Somesh Jha, Matt Fredrikson, Z Berkay Celik, and Ananthram Swami. 2016. The Limitations of Deep Learning in Adversarial Settings. In *Proceedings of the 1st IEEE European Symposium in Security and Privacy (EuroS&P)*.
- [28] Nicolas Papernot, Patrick McDaniel, Arunesh Sinha, and Michael Wellman. 2016. Towards the Science of Security and Privacy in Machine Learning. *arXiv preprint arXiv:1611.03814* (2016).
- [29] Nicolas Papernot, Patrick McDaniel, Xi Wu, Somesh Jha, and Ananthram Swami. 2015. Distillation as a Defense to Adversarial Perturbations against Deep Neural Networks. *Proceedings of the 37th IEEE Symposium on Security and Privacy (S&P)* (2015).
- [30] Saurabh Shintre Andrew B. Gardner Reuben Feinman, Ryan R. Curtin. 2017. Detecting Adversarial Samples from Artifacts. *CoRR abs/1703.00410* (2017). <http://arxiv.org/abs/1703.00410>
- [31] Paul R. Rosenbaum. An exact distribution-free test comparing two multivariate distributions based on adjacency. *Journal of the Royal Statistical Society B* (????), 2005.
- [32] Shiqi Shen, Shruti Tople, and Prateek Saxena. 2016. Auror: defending against poisoning attacks in collaborative deep learning systems. In *Proceedings of the 32nd Annual Conference on Computer Security Applications, ACSAC 2016, Los Angeles, CA, USA, December 5-9, 2016*. 508–519. <http://dl.acm.org/citation.cfm?id=2991125>
- [33] Nedim Srndic and Pavel Laskov. 2014. Practical Evasion of a Learning-Based Classifier: A Case Study. In *2014 IEEE Symposium on Security and Privacy, SP 2014, Berkeley, CA, USA, May 18-21, 2014*. 197–211. <https://doi.org/10.1109/SP.2014.20>
- [34] Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. 2014. Intriguing properties of neural networks. In *Proceedings of the 2014 International Conference on Learning Representations. Computational and Biological Learning Society*.
- [35] Gábor J. Székely and Maria L. Rizzo. 2013. Energy statistics: A class of statistics based on distances. *Journal of Statistical Planning and Inference* 143, 8 (2013), 1249 – 1272. <https://doi.org/10.1016/j.jspi.2013.03.018>
- [36] Beilun Wang, Ji Gao, and Yanjun Qi. 2016. A Theoretical Framework for Robustness of (Deep) Classifiers Under Adversarial Noise. *CoRR abs/1612.00334* (2016). <http://arxiv.org/abs/1612.00334>
- [37] Weilin Xu, Yanjun Qi, and David Evans. 2016. Automatically evading classifiers. In *Proceedings of the 2016 Network and Distributed Systems Symposium*.